

CSE258_HW4

March 6, 2017

1 CSE258 HW 4

```
In [1]: import numpy as np
import urllib
import string
import nltk
from nltk import bigrams
from collections import defaultdict
from sklearn import linear_model
from sklearn.metrics import mean_squared_error
from math import log10
from scipy import spatial
```

```
In [2]: def parseData(fname):
        for l in urllib.urlopen(fname):
            yield eval(l)

        print "Reading data..."
        data = list(parseData("http://jmcauley.ucsd.edu/cse190/data/beer/beer_50000"))
        print "done"
```

```
Reading data...
done
```

1.0.1 Task 1

There are 182246 unique bigrams among all of the reviews.

The 5 most-frequently-occurring bigrams along with their number of occurrences:

The bigram ('with', 'a') has an occurrence of 4587

The bigram ('in', 'the') has an occurrence of 2595

The bigram ('of', 'the') has an occurrence of 2245

The bigram ('is', 'a') has an occurrence of 2056

The bigram ('on', 'the') has an occurrence of 2033

```
In [3]: # bigram count
punctuation = set(string.punctuation)
bigramCount = defaultdict(int)
```

```

for d in data:
    text = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    bg = list(bigrams(text.split())) # all bigrams in text
    for b in bg:
        bigramCount[b] += 1
print 'There are ' + str(len(bigramCount)) + ' unique bigrams among all of the reviews'
# sort
bigramSorted = list(sorted(bigramCount, key = lambda x : bigramCount[x], reverse=True))
for i in range(0,5):
    print 'The bigram ' + str(bigramSorted[i]) + ' has an occurrence of ' + str(bigramCount[bigramSorted[i]])

```

There are 182246 unique bigrams among all of the reviews
The bigram ('with', 'a') has an occurrence of 4587
The bigram ('in', 'the') has an occurrence of 2595
The bigram ('of', 'the') has an occurrence of 2245
The bigram ('is', 'a') has an occurrence of 2056
The bigram ('on', 'the') has an occurrence of 2033

1.0.2 Task 2

The MSE obtained using bigrams only is : 0.34361068509441478

```

In [4]: Bigrams = [x for x in bigramSorted[:1000]]
        bigramID = dict(zip(Bigrams, range(len(Bigrams))))
        bigramSet = set(Bigrams)

def feat(data):
    feat = [0] * len(Bigrams)
    text = ''.join([c for c in data['review/text'].lower() if not c in punctuation])
    bg = list(bigrams(text.split())) # all bigrams in text
    for b in bg:
        if b in bigramSet:
            feat[bigramID[b]] += 1
    feat.append(1) #offset
    return feat

In [5]: x = [feat(d) for d in data]
        y = [d['review/overall'] for d in data]

        clf = linear_model.Ridge(1.0, fit_intercept=False)
        clf.fit(x, y)
        theta = clf.coef_
        predictions = clf.predict(x)

        mean_squared_error(predictions, y)

```

Out[5]: 0.34361068509441478

1.0.3 Task 3

The MSE obtained using unigrams and bigrams is : 0.28933386918744819

```
In [6]: # unigram count
wordCount = defaultdict(int)

for d in data:
    text = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    for w in text.split():
        wordCount[w] += 1

In [7]: # merge unigram and bigram
mergeCount = dict(wordCount)
mergeCount.update(bigramCount)

mergeSorted = list(sorted(mergeCount, key = lambda x : mergeCount[x], reverse=True))

grams = [x for x in mergeSorted[:1000]]
gramID = dict(zip(grams, range(len(grams))))
gramSet = set(grams)

def feat(data):
    feat = [0] * len(grams)
    text = ''.join([c for c in data['review/text'].lower() if not c in punctuation])
    bg = list(bigrams(text.split())) # all bigrams in text
    for w in text.split() + bg:
        if w in gramSet:
            feat[gramID[w]] += 1
    feat.append(1) #offset
    return feat

In [8]: x = [feat(d) for d in data]
y = [d['review/overall'] for d in data]

clf = linear_model.Ridge(1.0, fit_intercept=False)
clf.fit(x, y)
theta = clf.coef_
predictions = clf.predict(x)

mean_squared_error(predictions, y)
```

Out [8]: 0.28933386918744819

1.0.4 Task 4

The 5 unigrams/bigrams with the most positive associated weights are :

‘sort’ has associated weight of 0.521680776822

(‘a’, ‘bad’) has associated weight of 0.226288834348

('of', 'these') has associated weight of 0.22289001188
 ('not', 'bad') has associated weight of 0.216268615711
 ('the', 'best') has associated weight of 0.213772219036
 The 5 unigrams/bigrams with the most negative associated weights are :
 ('sort', 'of') has associated weight of -0.645937945334
 'water' has associated weight of -0.271900176951
 'corn' has associated weight of -0.23756003904
 ('the', 'background') has associated weight of -0.218138672448
 'straw' has associated weight of -0.199753548917

```

In [9]: weights = zip(theta[0:-1], range(1000))
        weights.sort()

        print 'The 5 unigrams/bigrams with the most positive associated weights are'
        for i in range(0,5):
            print str(grams[weights[-i-1][1]]) + ' has associated weight of ' + str(w
        print '\n The 5 unigrams/bigrams with the most negative associated weights'
        for i in range(0,5):
            print str(grams[weights[i][1]]) + ' has associated weight of ' + str(w
  
```

The 5 unigrams/bigrams with the most positive associated weights are :
 sort has associated weight of 0.521680776822
 ('a', 'bad') has associated weight of 0.226288834348
 ('of', 'these') has associated weight of 0.22289001188
 ('not', 'bad') has associated weight of 0.216268615711
 ('the', 'best') has associated weight of 0.213772219036

The 5 unigrams/bigrams with the most negative associated weights are :
 ('sort', 'of') has associated weight of -0.645937945334
 water has associated weight of -0.271900176951
 corn has associated weight of -0.23756003904
 ('the', 'background') has associated weight of -0.218138672448
 straw has associated weight of -0.199753548917

1.0.5 Task 5

The inverse document frequency of the words 'foam', 'smell', 'banana', 'lactic', and 'tart' are :

[1.1378686206869628, 0.5379016188648442, 1.6777807052660807, 2.9208187539523753,
 1.8068754016455384]

Their tf-idf scores in the first review (using log base 10) are:

[2.2757372413739256, 0.5379016188648442, 3.3555614105321614, 5.841637507904751,
 1.8068754016455384]

```

In [10]: wordList = ['foam', 'smell', 'banana', 'lactic', 'tart']
        N = len(data)
        # inverse document frequency
        def idf(w):
  
```

```

        count = 0
        for d in data:
            text = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
            if w in text.split():
                count += 1
        return log10(N * 1. / count)

idfList = [idf(w) for w in wordList]
print idfList

[1.1378686206869628, 0.5379016188648442, 1.6777807052660807, 2.9208187539523753, 1.1378686206869628]

```

```

In [11]: # term frequency in first review
def tf(w):
    count = 0
    text = ''.join([c for c in data[0]['review/text'].lower() if not c in punctuation])
    for t in text.split():
        if t == w:
            count += 1
    return count

tfList = [tf(w) for w in wordList]

In [12]: [i * j for i, j in zip(idfList, tfList)]

Out[12]: [2.2757372413739256,
          0.5379016188648442,
          3.3555614105321614,
          5.841637507904751,
          1.8068754016455384]

```

1.0.6 Task 6

The cosine similarity between the first and the second review in terms of their tf-idf representations is: 0.10613024167865803

```

In [13]: counts = [(wordCount[w], w) for w in wordCount]
counts.sort()
counts.reverse()

# tfidf feature words -- 1000 most common
featWords = [x[1] for x in counts[:1000]]
featID = dict(zip(featWords, range(1000)))

# idf for feature words
countList = [0] * 1000
for d in data:
    text = ''.join([c for c in d['review/text'].lower() if not c in punctuation])

```

```

        textWord = set(text.split())
        for t in textWord:
            if t in featID:
                countList[featID[t]] += 1
idfList = [log10(N * 1. / x) for x in countList]

In [14]: # tfidf feature extraction
def tfidfFeat(d):
    count = [0] * len(featWords) # tf counts
    text = ''.join([c for c in d['review/text'].lower() if not c in punctu
    for t in text.split():
        if t in featWords:
            count[featID[t]] += 1
    feat = [i * j for i, j in zip(idfList, count)]
    # feat.append(1)
    return feat

In [15]: feat1 = tfidfFeat(data[0])
        feat2 = tfidfFeat(data[1])
        1 - spatial.distance.cosine(feat1, feat2)

Out[15]: 0.10613024167865803

```

1.0.7 Task 7

beerID : 52211

profileName : Dope

reviewText : A: A hazy deep orange pour, almost red. Small white head that fades quickly. A little spotty lacing. S: Big pumpkin, cinnamon, ginger, nutmeg and brown sugar. Sweet. Smells like a pumpkin pie mixed with a gingerbread cookie. T: Tons of pumpkin dominates throughout. Cinnamon, ginger, nutmeg and a bit of vanilla creaminess. M: Smooth medium body. Tiny bit of drying alcohol. O: Excellent pumpkin ale. Heavy on the pumpkin but the spices don't get completely overshadowed either.

```

In [17]: cosineList = [1 - spatial.distance.cosine(feat1, tfidfFeat(x)) for x in da
        result = zip(cosineList, range(len(cosineList)))
        result.sort()
        result.reverse()
        print 'beerID : ' + str(data[result[0][1]]['beer/beerId']) + '\n' + 'profi
        print 'reviewText : ' + data[result[0][1]]['review/text']

```

beerID : 52211

profileName : Dope

reviewText : A: A hazy deep orange pour, almost red. Small white head that fades qu

1.0.8 Task 8

The MSE obtained with the 1000-dimensional tf-idf representations is : 0.27875956007772285

```
In [18]: # rewrite the tfidf feature extraction with offset
def tfidfFeat(d):
    count = [0] * len(featsWords) # tf counts
    text = ''.join([c for c in d['review/text'].lower() if not c in punctu
    for t in text.split():
        if t in featsWords:
            count[featsID[t]] += 1
    feat = [i * j for i, j in zip(idfList, count)]
    feat.append(1)
    return feat

In [19]: x = [tfidfFeat(d) for d in data]
y = [d['review/overall'] for d in data]

clf = linear_model.Ridge(1.0, fit_intercept=False)
clf.fit(x, y)
theta = clf.coef_
predictions = clf.predict(x)

mean_squared_error(predictions, y)

Out[19]: 0.27875956007772285

In [ ]:
```